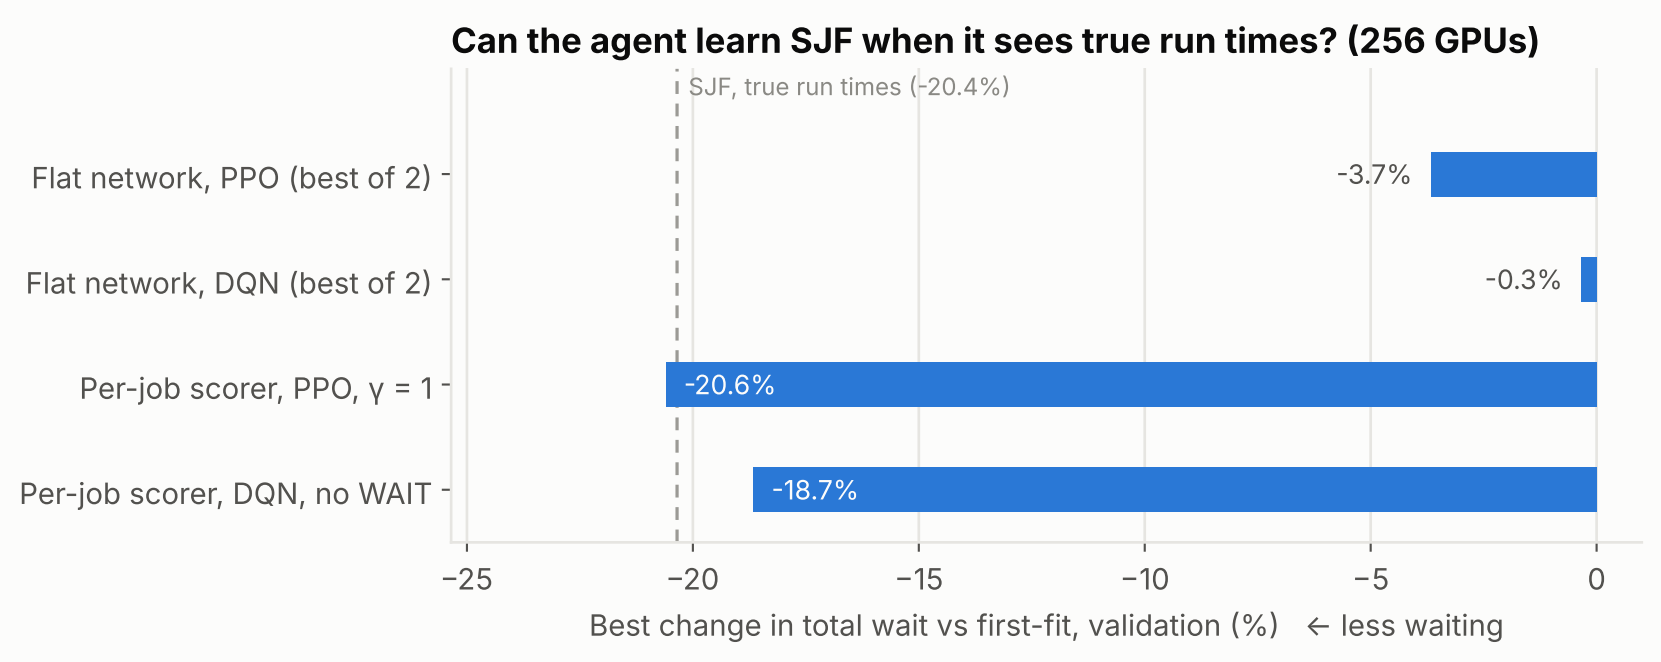
Can the agent learn SJF when it sees true run times? (256 GPUs)
SJF, true run times (-20.4%)
Flat network, PPO (best of 2)
-3.7%
Flat network, DQN (best of 2)
-0.3%
Per-job scorer, PPO, γ = 1
-20.6%
Per-job scorer, DQN, no WAIT
-18.7%
−25
−20
−15
−10
−5
0
Best change in total wait vs first-fit, validation (%) ← less waiting